

# COMP 543: Tools & Models for Data Science

## Relational Algebra

Chris Jermaine & Risa Myers

Rice University



# Relational Calculus vs. Algebra

- In Relational Calculus
  - You say what you want
  - And not how to compute it
- But obviously...
  - This needs to be compiled into an actual computational plan
  - And in relational DBs, the plan is expressed in relational algebra
- RA is the “abstract machine” of relational databases

# What Is An Algebra?

- Many Definitions!
  - Simplest: it is a set (domain) with a number of operations
  - The domain is closed under those operations
- In RA...
  - The domain is the set of all valid relations
  - The set of operations includes  $\pi, \sigma, \times, \bowtie, \cup, \cap, -$
- Now let's go through the operations!

# Projection

- Projection removes attributes
- $\pi_A(R)$ ...
  - $A$  is a set of attributes of relation  $R$
  - This simply removes all attributes not in  $A$  from  $R$
  - Note: cardinality of output can differ from  $R$
  - Output is a relation

FREQUENTS

DRINKER	CAFE
Risa	JL
Risa	BH
Chris	BH
Chris	DT

$\pi_{\text{DRINKER}}(\text{FREQUENTS})$

DRINKER
Risa
Chris

# Selection

- Selection removes tuples
- $\sigma_B(R)$ ...
  - $B$  is a boolean predicate that can be applied to a single tuple from  $R$
  - This simply removes all tuples not accepted by  $B$
  - Again: output is a relation

FREQUENTS

DRINKER	CAFE
Risa	JL
Risa	BH
Chris	BH
Chris	DT

$\sigma_{\text{DRINKER}='Risa'}(\text{FREQUENTS})$

DRINKER	CAFE
Risa	JL
Risa	BH

# Selection/Projection Example

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

? Query: Who likes 'Cold Brew' coffee?

# Selection/Projection Example

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Query: Who likes 'Cold Brew' coffee?
  - $\pi_{\text{DRINKER}}(\sigma_{\text{COFFEE}=\text{'Cold Brew'}}(\text{LIKES}))$

# Join: Cartesian Product

- Join combines tuples
- Simplest join is Cartesian product (aka: cross product)
- $R \times S \dots$ 
  - Returns  $r \bullet s$  for all  $r \in R, s \in S$
  - ? What is the output cardinality?



LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Often you want  $\sigma_B(R \times S)$
- Shorthand for this is  $R \bowtie_B S$
- ? Query: Who likes a coffee that 'Risa' likes?

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Often you want  $\sigma_B(R \times S)$
- Shorthand for this is  $R \bowtie_B S$
- Query: Who likes a coffee that 'Risa' likes?
  - $\text{TEMP}(d_1, c_1, d_2, c_2)$   
 $\leftarrow \text{LIKES} \bowtie_{\text{COFFEE}=\text{COFFEE}} (\sigma_{\text{DRINKER}=\text{'Risa'}}(\text{LIKES}))$
  - $\pi_{d_1}(\text{TEMP})$

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Often you want to join two relations
  - Using an equality check on all attributes having the same name
  - Then project away redundant attributes
- Shorthand for this is  $R * S$
- ? Query: Who goes to a cafe serving a coffee that they like?

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Often you want to join two relations
  - Using an equality check on all attributes having the same name
  - Then project away redundant attributes
- Shorthand for this is  $R * S$
- Query: Who goes to a cafe serving a coffee that they like?
  - $\pi_{\text{DRINKER}}(\text{LIKES} * \text{FREQUENTS} * \text{SERVES})$

# Set-Based Operations

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Can use standard set operations as well:  $\cup, \cap, -$ 
  - To use, types and numbers of input attributes must match
  - By convention, attribute names come from LHS
  - $R \cup S$ : all tuples in  $R$  or in  $S$
  - $R \cap S$ : all tuples in  $R$  and in  $S$
  - $R - S$ : all tuples in  $R$  and not in  $S$

? Query: Who does not like 'Cold Brew' coffee?

# Set-Based Operations

LIKES (DRINKER, COFFEE)  
FREQUENTS (DRINKER, CAFE)  
SERVES (CAFE, COFFEE)

- Can use standard set operations as well:  $\cup, \cap, -$ 
  - To use, types and numbers of input attributes must match
  - By convention, attribute names come from LHS
  - $R \cup S$ : all tuples in  $R$  or in  $S$
  - $R \cap S$ : all tuples in  $R$  and in  $S$
  - $R - S$ : all tuples in  $R$  and not in  $S$
- Query: Who does not like 'Cold Brew' coffee?
  - $\text{COFFEEGOOD} \leftarrow \pi_{\text{DRINKER}}(\sigma_{\text{COFFEE}='Cold Brew'}(\text{LIKES}))$
  - $(\pi_{\text{DRINKER}}(\text{LIKES})) - \text{COFFEEGOOD}$

# Complicated Set-Based Example

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

? Who only goes to cafes where they can get a coffee they like?

# Complicated Set-Based Example

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Who only goes to cafes where they can get a coffee they like?
  - Use 'all people' – 'those who go to a cafe where they can't get a coffee they like'
  - $ALLPEEPS \leftarrow \pi_{DRINKER}(LIKES)$
  - How about 'those who go to a cafe where they can't get a coffee they like'?
  - Use FREQUENTS – 'DRINKER, CAFE combos where the person can get a coffee they like'
  - $GOODCOFFEE \leftarrow \pi_{DRINKER,CAFE}(LIKES * SERVES)$
- Then the answer is
  - $ALLPEEPS - \pi_{DRINKER}(FREQUENTS - GOODCOFFEE)$



# Questions?